

# Parallelization in OpenFOAM for HPC Deployment

Hands-on activities

---

Mohammed Elwardi Fadel<sup>1,2</sup>, Holger Marschall<sup>1</sup> and Christian Hasse<sup>2</sup>

April, 2023

<sup>1</sup> Mathematical Modeling and Analysis (MMA)

<sup>2</sup> Simulation of Reactive Thermo Fluid Systems (STFS)  
Energy Conversion Group, NHR4CES - TU Darmstadt

## Introductory activities

## Activity 01: Blocking P2P comms - A first look

- `parallelClass` class has a `lst_` member, initiated differently on each process.
- Your task is to append `lst_` from all **slave** processes (1, 2 and 3) into a single list object on master.
  - Using blocking P2P communication between master and a slave each time.
  - Append the lists in the order of process IDs.
  - Modify `exercises/parallelClass.C` file.




Figure 1: Blocking P2P with master

## Activity 02: Collective comms - A first look

- `parallelClass::isPrime(int)` is a primitive method to check if its argument is a prime number.
  - But it's serial code! Running it in parallel will **duplicate work** on all processes
  - Parallelizing it should yield performance gains

### 1. Domain Decomposition

- This is done already if your data is mesh-related (mesh itself, fields, ... etc)
- We need to "decompose" the search range `[3, sqrt(n)+1]` into `nProcs` ranges.
- `parallelClass::next(int n, int& i)` increments `i` to the next number in process-controlled range



Figure 2: Provided trivial decomposition of the range  $[3, \sqrt{n} + 1]$

## Activity 02: Collective comms - A first look

### 2. Blocking P2P comms:

- Transform `parallelClass::isPrime(int)` so that it acts on the corresponding range on each of the 4 processes.
- All processes **must** decide if it's a prime number, based on results from all other process.
- Think: collective comms. Although P2P comms also would work!

### 3. An alternate decomposition:

- Each process is responsible for possible divisors that are `nProcs` apart.
- Implement this behavior in `parallelClass::next(int n, int& i)`



Figure 3: A new example decomposition of the range  $[3, \sqrt{n} + 1]$

## Activity 02: Collective comms - A first look

Is there room for improvement?

### 4. Decomposition effect


- Different decompositions result in different loads on processes
- In blocking comms, balanced decompositions are the most efficient
- But keep in mind; if it takes too much time to decompose "dynamically" - it might backfire!

### 5. Optimizing the code

- A process can **break** out of loop and return if it finds a viable divisor.
- But can't stop calculations on all other processes.
- Hence, apparently, not worth the trouble.
- **But**, if all other processes also return early, we can gain some CPU time.
- Maybe test only on **prime numbers** less than  $\sqrt{n} + 1$ ?
- The whole parallelization idea was premature optimization, huh!

## Activity 03: Blocking P2P comms are not good for your health

- Code in `parallelClass::run()` tries to perform 2 blocking P2P communication operations between two processes.
  - When they hang; processes are killed by `timeout`.
- Your task is to modify the `run` member method so the code no longer hangs.
  - Pay attention to the order of the send/receive ops!
  - The Actual sending happens on `*Pstream` objects' destruction
  - Do you think your solution would withstand MPI implementation swapping? (OpenMPI, Intel MPI, MPICH ... etc)



**Figure 4:** 2 Blocking P2P communication operations at the same time

## Activity 04: Non-Blocking P2P comms for swapping ops

- Each process communicates a list object to its **neighbour**
  - Neighbouring relationships are based on mesh decomposition
  - **parallelClass** class has **nProcs** lists as a member variable **lists\_** (Each list is to be transferred to the process's neighbour)
- Your task is to perform an exchange of the lists so that, at the end of all transfers; each process holds the lists from its neighbors (instead of its own ones).
  - This should be done in **parallelClass::swapLists()** method
  - Again in **execises/parallelClass.C** file



## Activity 04: Non-Blocking P2P comms for swapping ops





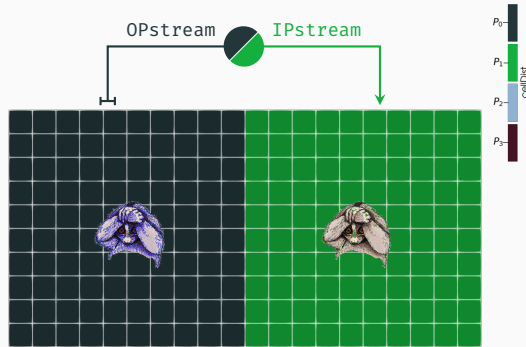
Figure 5: Default (hierarchical) decomposition of the block-mesh in the cavity case

## Activity 04: Non-Blocking P2P comms for swapping ops



**Figure 5:** Default (hierarchical) decomposition of the block-mesh in the cavity case

## Activity 04: Non-Blocking P2P comms for swapping ops



**Figure 5:** Default (hierarchical) decomposition of the block-mesh in the cavity case

## Activity 05: Collective comms - Reference cell

- It's popular to assign a **reference cell** if all BCs are gradient-based for a field:

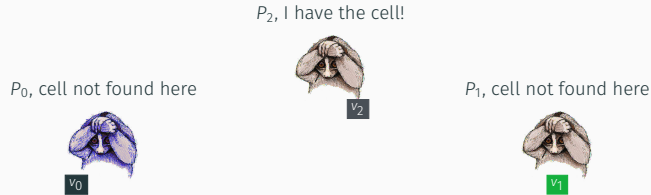
Common pressure reference point in `system/fvSolution`

```
1 PISO
2 {
3     pRefPoint    (0.0975 0.0025 0);           // At far right-bottom cell of the cavity mesh
4     pRefValue    0;
5 }
```

- Note that `pRefCell` is enough for most cases. Information propagated automatically through processor patches!
- `mesh.findCell(position);` returns:
    - `-1` if the position is outside the local mesh (on calling process).
    - `cellID` of the cell containing the position otherwise.
1. Your first task is to check if the provided position is inside the **global mesh**
    - Using collective comms; Think: `Foam::reduce`
    - Change `parallelClass::checkPosition(const vector&)` so it returns `true` if given position is inside the mesh.

## Activity 05: Collective comms - Reference cell

2. It's useful for everyone to know **which process has the corresponding reference cell**
- Change `parallelClass::whoHasReferenceCell(const vector&)` so it returns -1 if position is outside the mesh; and returns the rank of the process which is responsible for the reference cell otherwise.
  - Again, using collective comms.



**Figure 6:** Suggested setup for reference cell communication

## Activity 05: Collective comms - Reference cell

2. It's useful for everyone to know **which process has the corresponding reference cell**
- Change `parallelClass::whoHasReferenceCell(const vector&)` so it returns `-1` if position is outside the mesh; and returns the rank of the process which is responsible for the reference cell otherwise.
  - Again, using collective comms.




Figure 6: Suggested setup for reference cell communication

## Activity 05: Collective comms - Reference cell

- It's useful for everyone to know **which process has the corresponding reference cell**
  - Change `parallelClass::whoHasReferenceCell(const vector&)` so it returns `-1` if position is outside the mesh; and returns the rank of the process which is responsible for the reference cell otherwise.
  - Again, using collective comms.




Figure 6: Suggested setup for reference cell communication


## Activity 06: Parallel comms for custom data types

- So, I wrote a class; now I want its objects to travel through processors.
  - `parallelClass.H` file declares an `Edge` class, to represent "neighbouring relationship" between two processors.
  - The goal is to build a graph of such edges.
- 1. Your first task is to make this `Edge` class compatible with random-access lists. I.e an `Edge` object can be put in a list.
  - Try `./Alltest` and take a look at compilation log.
  - Modify `Edge` class declaration/definition so that errors related to its compatibility with `List` template disappear.



## Activity 06: Parallel comms for custom data types


2. Implement **operator<<** and **operator>>** so that an **Edge** object can be passed-to/read-from Output/Input streams.
  - Things should at least compile.
  - Compatibility of **Edge** with **List** is tested by **the compiler**.
  - There is also a check for correct graph communications when the graph is gathered then scattered.



**Figure 7:** Typical way of sending custom objects through **\*Pstreams** using Collective Comms

## Activity 07: Special parallel comms for custom data types


- All good, but my class can't have a null constructor!
  - Best solution is to use **Linked Lists**.
  - If constructor arguments are needed on the other end; you need a factory sub-class.
- **Edge** class was modified to fit into linked lists, and can be constructed from a mesh instance and an input stream.
  - Your task is to make it **\*Pstreams-ready**



**Figure 8:** Typical way of sending custom objects through **\*Pstreams** using P2P Comms

## Activity 07: Special parallel comms for custom data types

- All good, but my class can't have a null constructor!
  - Best solution is to use **Linked Lists**.
  - If constructor arguments are needed on the other end; you need a factory sub-class.
- **Edge** class was modified to fit into linked lists, and can be constructed from a mesh instance and an input stream.
  - Your task is to make it **\*Pstreams**-ready



**Figure 8:** Typical way of sending custom objects through **\*Pstreams** using P2P Comms

Final projects

# Project 01: Parallelizing a coded source term

→ This project does not work with Foam-Extend as it doesn't support `fvOptions`.


1. You inherit a repository with a cavity case which works "as expected" in serial
2. Run `./Allrun` from inside the case's directory to compare results from serial and parallel runs.
  - This requires ParaView to be installed.
  - Use the server version for headless machines if you're on a container.
  - The script writes a `log.pvpython` file containing max/min absolute error in velocity values at  $t = 0.5s$  between serial and parallel runs of a cavity case.
3. The provided cavity case has a coded `fvOptions` which adds a source term to the velocity equation.
  - Run `./Allrun` with the source active
  - Disable the `fvOptions` source and `./Allclean && ./Allrun` again.
    - By either setting `codedSource.active` to false, in `system/fvOptions`.
    - Or by moving `fvOptions` file elsewhere.

# Project 01: Parallelizing a coded source term

How does the custom source work?

- We define a box inside our cavity-case mesh to add an explicit vector source to the **UEqn** there.
- The x-component of the source at each cell  $S_i = S_v * \frac{1}{nNeighbours+1} (\sum_{j=1}^{nNeighbours} k_j + k_i)$  depends on:
  1. Some "total" source value  $S_v$  provided by the user.
  2. An average value of a coefficient field  $k$  over the cell and its immediate neighbours.
- the target volume is defined through spacial dimensions and the corresponding cells are found through a Cell Set: **box (0.02 0.02 -1) (0.06 0.06 1)**

## Project 01: Parallelizing a coded source term



**Figure 9:** A sample run of the randomized coefficient field  $k$ . (The white box is where the source is applied)

# Project 01: Parallelizing a coded source term

## Suggested Steps

1. Identify the problem lines in fvOptions' code.

Hint: Look for lines which use local mesh information!

2. Implement a fix for the identified issues.

Hint: Can we get information about neighbouring cells that are on the other processor?

3. How optimized/sophisticated do you think your solution is? Share it with your peers and take a look at theirs!



## Project 01: Parallelizing a coded source term

To help you identify the issue, take a look at this parameter variation study (varying source box dimensions,  $S_v$  parameter and number of MPI processes involved):

**Table 1:** Sample trial data for the parameter variation study on the cavity case

Trial	$S_v$	$x_{min}$	$x_{max}$	$y_{min}$	$y_{max}$	nProcs	MaxError
0	9.569e-4	0.01	0.07	0.04	0.06	5	0.75085
1	8.102e-4	0.01	0.06	0.03	0.06	8	0.36281
2	3.061e-5	0.03	0.06	0.04	0.06	7	0.01018
3	8.598e-4	0.03	0.06	0.01	0.05	4	0.65839

# Project 01: Parallelizing a coded source term




Figure 10:  $k$  coefficient field for a sample from trial data

# Project 01: Parallelizing a coded source term

Want a quick way to test your fix?

Conduct your own parameter variation studies while you fix the case!

```
1  # Clone the helper repository
2  git clone https://github.com/FoamScience/OpenFOAM-Multi-Objective-Optimization multiOptFoam
3  cd multiOptFoam
4  # Install dependencies
5  pip3 install -r requirements.txt
6  # Copy your case
7  cp -r /path/to/cavity/case template
8  # Grab the config file (provided with the case)
9  cp /path/to/config.yaml .
10 # Run parameter variation
11 ./paramVariation.py
12 # Now change template/system/fvOptions, clean and rerun the variation study
13 rm -rf Example* && ./paramVariation
14 # The goal is to get the maxError column for all trials as close to zero as possible (~ 1e-6)
```